

CS/ECE 4457

Computer Networks: Architecture and Protocols

Lecture 4/5

- Three Architectural Principles
 - Design Goals

Qizhe Cai



Announcements

- **Exam conflict:**
 - Today is the last day to announce your exam conflicts.
 - Exam 1: 02/10
 - Exam 2: 03/05
 - Exam 3: 04/28
 - Thank you to those who already sent us an email
 - We will send an email to all those who have a conflict
- Problem set 1 solutions released (on Piazza)
- Problem set 2 will be released soon (on course website)

Context for Today's Lecture

- So far, we have discussed several high-level concepts
 - Network sharing
 - End-to-end working of the Internet
 - Addressing, Routing, Switch/Router functionality, etc.
- And, have dived deep into several topics:
 - Circuit switching and packet switching (especially the “why”)
 - Delays (transmission, propagation)
- **You know more about computer networks than you may realize!**
- **Today: Continue to lay the foundation for rest of the course**

Goals for Today's and next Lecture

- Three architectural principles:
 - Layering
 - End-to-end principle
 - Fate Sharing principle
- Design goals for computer networks:
 - Eight of them
- We will come back to these over and over again
 - Almost every lecture in the semester
- Before we start, let me outrightly admit
 - First time I learnt these, I said — what the @#\$%
 - ... there are easier ways to torture students!
 - Now, these have become the guiding principles of my career!

Quick recap from last lecture

Recap: four fundamental problems!

- **Locating the destination:** Naming, addressing
 - Mapping of names to addresses using Domain Name System
- **Finding a path to the destination:** Routing
 - Distributed algorithm that computes and stores routing tables
- **Sending data to the destination:** Forwarding
 - Input queues, virtual output queues, output queues
 - Enablers: Packet header (address), and routing table (outgoing link)
- **Reliability:** Failure handling
 - Not much discussion, but **the** question: hosts or networks?

Recap: the final piece in the story – Host network stack

Of Sockets and Ports

- When a process wants access to the network, it opens a socket, which is associated with a port
- **Socket:** an OS mechanism that connects processes to the network stack
- **Port:** number that identifies that particular socket
- The port number is used by the OS to direct incoming packets

Recap: Implications for Packet Header

- **Packet Header must include:**
 - Source and destination address (used by network)
 - Source and destination port (used by network stack)
- When a packet arrives at the destination host, packet is delivered to the socket associated with the destination port
- More details later

Recap: the end-to-end story

- Application opens a **socket** that allows it to connect to the **network stack**
- Maps **name** of the web site to its **address** using **DNS**
- The network stack at the source embeds the address and **port** for both the source and the destination in **packet header**
- Each **router** constructs a **routing table** using a distributed algorithm
- Each router uses destination address in the packet header to look up the **outgoing link** in the routing table
 - And when the link is free, forwards the packet
- When a packet arrives the destination:
 - The network stack at the destination uses the port to forward the packet to the right application

Recap: Separation of concerns

- **Network fabric:** Deliver packets from stack to stack (based on address)
- **Network stack (OS):** Deliver packets to appropriate socket (based on port)
- **Applications:**
 - Send and receive packets
 - Understand content of packet bodies

Questions?

Who cares?

- Why is separation of concerns important?
 - Separation of concerns ~ Modularity
- If each component's task well-defined, one can focus design on that task
 - And replace it with any other implementation that does that task
 - Without changing anything else

What is Modularity

- Modularity is nothing more than decomposing programs/systems into smaller units.
 - **A clean “separation of concerns”**
- Plays a crucial role in computer science...
- ... and networking

Computer System Modularity

- Partition system into modules
 - Each module has well defined interface
- Interfaces give flexibility in implementation
 - Changes have limited scope
- Examples
 - Libraries encapsulating set of functionalities
- The trick is to find the *right* modularity
 - The interfaces should be long-lasting
 - If interfaces are changing often, modularity is wrong

Network System Modularity

- The need for modularity still applies
 - **And is even more important! (why?)**
- Network implementations not just distributed across many lines of code
 - Normal modularity “organizes” that code
- Networking is distributed across many machines
 - Hosts
 - Routers

“Thinking” Network System Modularity

- Applications deal with data
- End-host network **stacks** move data from applications to the fabric
- Network **fabric** delivers data between **network stacks**
- **Network (stack + fabric)** delivers data **between applications**
- What is the **interface** between applications and network stacks?
 - **Sockets**
- What is the **interface** between network stacks and network fabric?
 - **Packet headers**
- The right way to think about sockets and packets

Three Architectural Principles

Network Modularity Decisions

- How to break system into modules?
 - Classic decomposition into tasks
- Where are modules implemented?
 - Hosts?
 - Routers?
 - Both?
- Where is state stored?
 - Hosts?
 - Routers?
 - Both?

Leads to three design principles

- How to break system into modules
 - **Layering**
- Where are modules implemented
 - **End-to-End Principle**
- Where is state stored?
 - **Fate-Sharing**

Layering

Breakdown end-to-end functionality into tasks

- Bits on wire
- Packets on wire
- Deliver packets between hosts in a “local” network (eg, within UVA)
- Routing & forwarding packets across networks (eg, from UVA to UIUC)
- Deliver data reliably between processes (applications)
- Do something with the data

Breakdown end-to-end functionality into tasks

- Bits on wire
- Packets on wire
- **Deliver packets between hosts in a local network**
- **Routing and forwarding (packets) across networks**
- **Deliver data reliably between processes**
- Do something with the data

Resulting Modules (Layers)

- Bits on wire (Physical)
- Packets on wire
- Deliver packets between hosts in a local network (Datalink)
- Routing and forwarding (packets) across networks (Network)
- Deliver data reliably between processes (Transport)
- Do something with the data (Application)

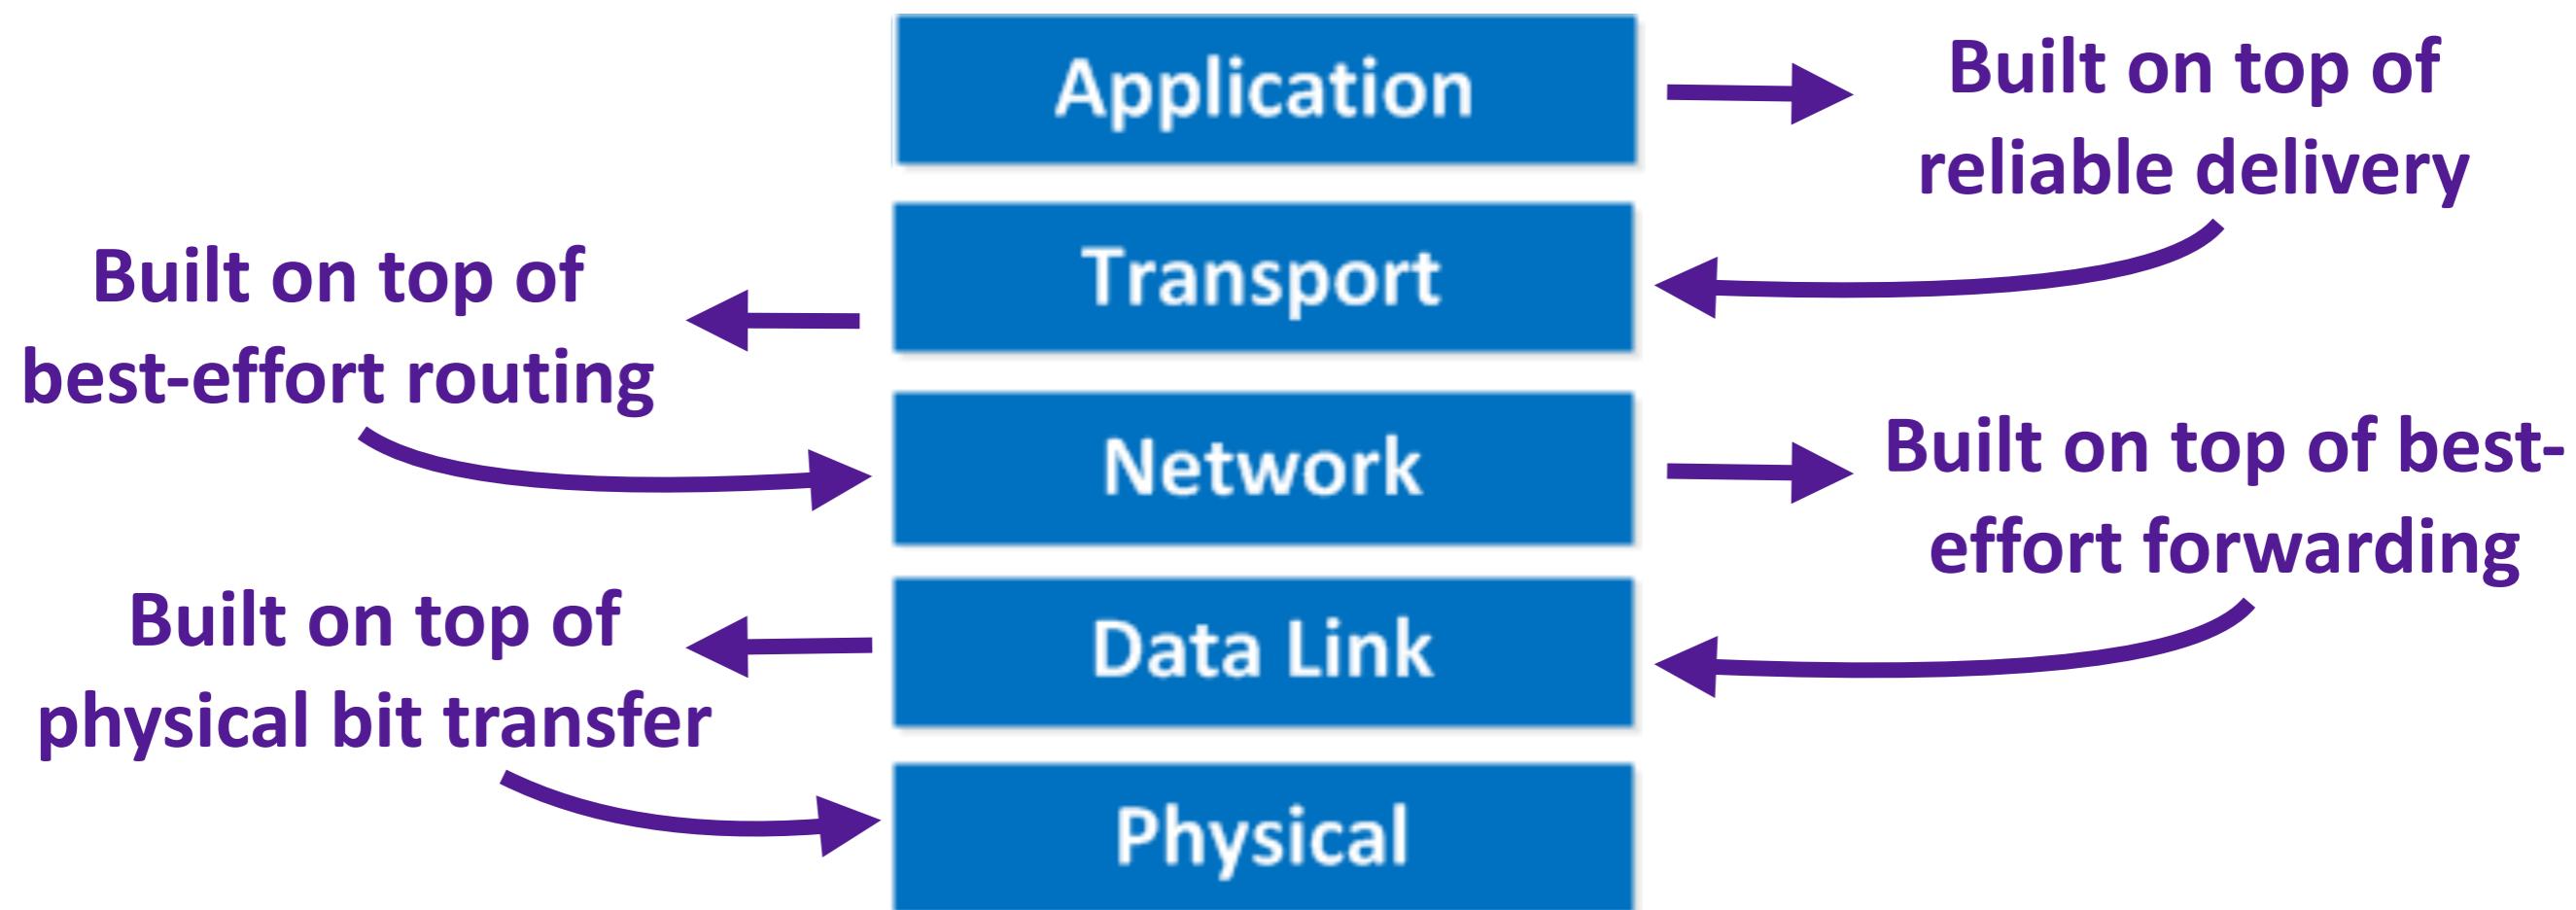
Resulting Modules (Layers)

- Bits on wire (Physical, Layer1)
- Packets on wire
- Deliver packets to hosts across local network (Datalink, Layer2)
- Routing and forwarding (packets) across networks (Network, Layer3)
- Deliver data reliably between processes (Transport, Layer4)
- Do something with the data (Application)

Five Layers (Top - Down)

- Application: Providing network support for apps
- **Transport (L4):** (Reliable) end-to-end delivery
- **Network (L3):** Routing and forwarding across networks
- **Datalink (L2):** Forwarding within a local network
- **Physical (L1):** Bits on wire

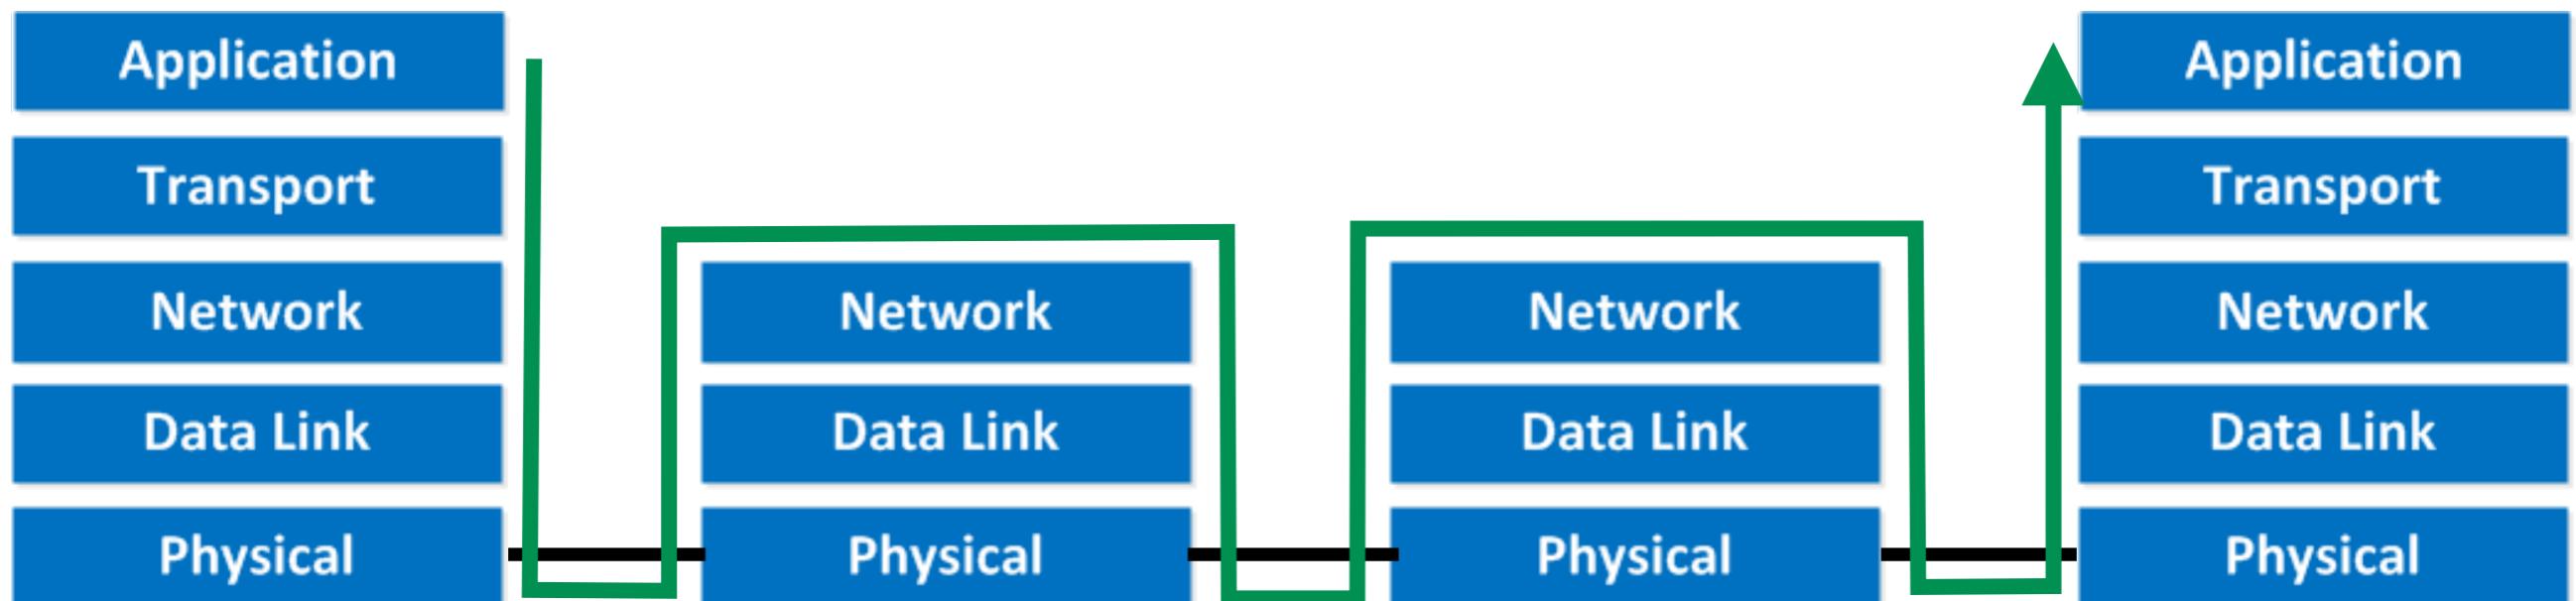
Layering



- A kind of modularity
 - Functionality separated into layers
 - Layer n **interfaces with only layer n-1 and layer n+1**
 - Hides complexity of surrounding layers

An end-to-end view of the layers

- Application: Providing network support for apps
- **Transport (L4)**: (Reliable) end-to-end delivery
- **Network (L3)**: Routing and forwarding across networks
- **Datalink (L2)**: Forwarding within a local network
- **Physical (L1)**: Bits on wire



Why does the packet go all the way to network layer at each hop?

Questions?

Three Internet Design Principles

- How to break system into modules?
 - Layering
- Where are modules implemented?
 - **End-to-End Principle**
- Where is state stored?
 - Fate-Sharing

Distributing Layers across Network

- Layers are simple if only on a single machine
 - Just stack of modules interacting with those above/below
- But we need to implement layers across machines
 - Hosts
 - Routers/switches
- What gets implemented where? And why?

What gets implemented on Host?


- Bits arrive on wire, must make it up to application
- Therefore, all layers must exist at host!

What gets implemented on Router?

- Bits arrive on wire
 - Physical layer necessary
- Packets must be forwarded to next router/switch
 - Datalink layer necessary
- Routers participate in global delivery
 - Network layer necessary
- **Routers do not support reliable delivery**
 - Transport layer (and above) not supported
 - Why?

Visualizing what gets implemented where

- Lower three layers implemented everywhere
- Top two layers only implemented at hosts



But why implemented this way?

- Layering tells you **what services each layer should provide**
- But doesn't tell you which layer should be implemented on which nodes


End-to-end Principle

If a function can completely and correctly be implemented only with the knowledge and help of the application standing at the endpoints of the communication system,

then providing that function as a feature of the communication system itself is not possible.

Sometimes providing an incomplete version of that function as a feature of the communication system itself may be useful as a performance enhancement.

End-to-end Principle: an example



- Suppose each link layer transmission is reliable
 - Does that ensure end-to-end (application-to-application) reliability?
- Suppose network layer is reliable
 - Does that ensure end-to-end (application-to-application) reliability?

End-to-end Principle: lets read again

If a function can completely and correctly be implemented only with the knowledge and help of the application standing at the endpoints of the communication system,

then providing that function as a feature of the communication system itself is not possible.

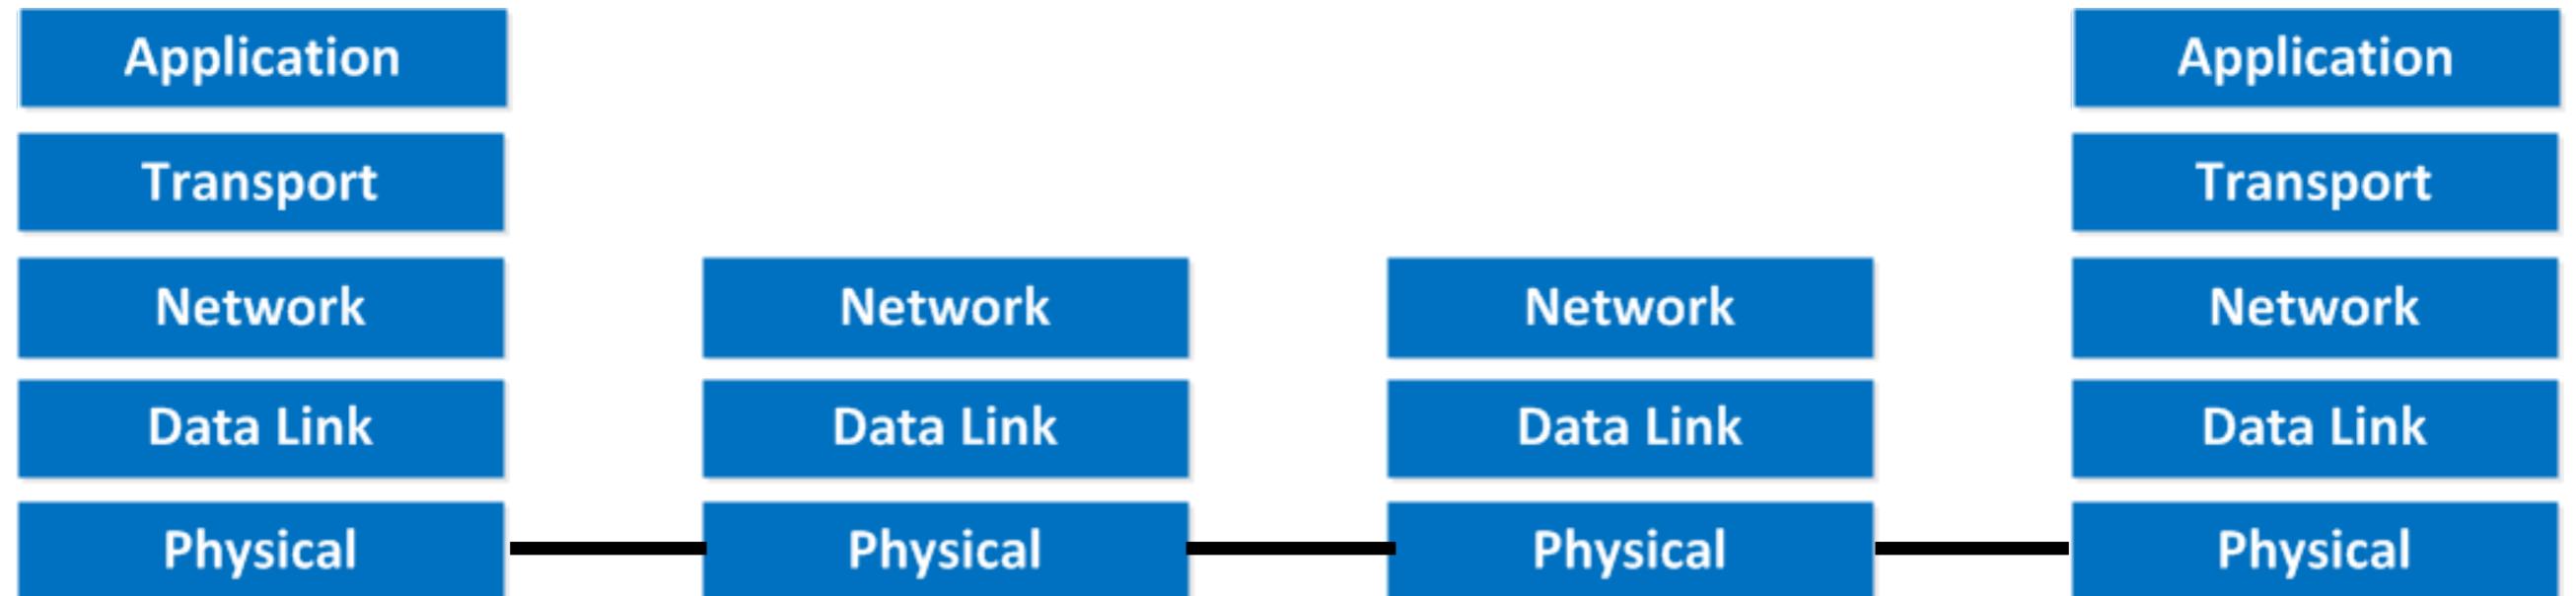
Sometimes providing an incomplete version of that function as a feature of the communication system itself may be useful as a performance enhancement.

End-to-end Principle (Interpretation)

Assume the condition (IF) holds. Then,

- **End-to-end implementation**
 - Correct
 - Generalized, and simplifies lower layers
- **In-network implementation**
 - Insufficient
 - May help — or hurt — performance

End-to-end Principle (Interpretation)



What does the end mean?

End-to-end Principle (Three things to know)

- **Everyone knows what it is**
 - So, you must!
- **Everyone believes it**
 - So, you must!
- **Nobody knows what it means**
 - We are all doomed anyways.

Questions?

Three Internet Design Principles

- How to break system into modules?
 - Layering
- Where are modules implemented?
 - End-to-End Principle
- **Where is the state stored?**
 - Fate-sharing

General Principle: Fate-Sharing

- When storing state in a distributed system, colocate it with entities that rely on that state
 - e.g, Connection states vs. Transport layers
- Only way failure can cause loss of the critical state is if the entity that cares about it also fails ...
 - ... in which case it doesn't matter
- Often argues for keeping network state at end hosts rather than inside routers
 - E.g., packet switching rather than circuit switching

Questions?

Decisions and their Principles

- How to break system into modules
 - **Dictated by layering**
- Where modules are implemented
 - **Dictated by End-to-End Principle**
- Where state is stored
 - **Dictated by Fate Sharing**